

Arquitetura de Computadores 2023/24

TPC4

This homework consists of a programming exercise to be completed **in a group of no more than two students**. You can discuss general doubts with colleagues, but the solution and the writing of the code must be strictly carried out by the members of the group. All solutions will be automatically compared, and cases of plagiarism will be punished according to current regulations. **If you use tools such as CoPilot or ChatGPT, you must include a comment in the source code reporting such use.**

Deadline for submission: 28/5 (Tuesday) at 10:00.

Dynamic memory allocation and mmap OS call

During the execution of a program, obtaining space in RAM is typically based on the C library functions malloc and free. The memory space used for this is called the heap, which is a continuous space that can be altered with the sbrk system call.

The following figure summarizes the memory map of a process in an operating system like Linux:




Figure 1 - Memory map of a Linux process

A process can obtain space in RAM through the mmap system call. This system call has a complex set of parameters that can be consulted using the command man mmap, but in the context of this work, we assume the set of parameters exemplified in the following C program:

```
#include <sys/mman.h>

size_t size = ... ; // size must be a multiple of the size of a page (4096)

void *addr = mmap( 0, size, PROT_READ | PROT_WRITE , MAP_PRIVATE | MAP_ANONYMOUS, -1, 0 );
if( addr == (void *)-1 ){
    perror("mmap");
    return 1;
}
```

If the system call is successful, the memory map becomes:




Figure 2 - Memory map after using the `mmap` system call

The process has access to the new area via C instructions such as:

```
char *pt = addr;

pt[5] = pt[22] + 5; // reading and writing in the new allocated memory area
```

In this work, we will implement a set of operations (API) that allow the use of this memory area through operations similar to the `malloc` and `free` operations in the C library. The API is defined in the `myAlloc.h` file whose content is as follows:

```
typedef struct {
    void *base;      // beginning of managed area
    void *top;       // address of last byte allocate
    void *limit;     // last address in the area allocated by mmap
} heap;

typedef struct{
    unsigned long long int available;
    unsigned long long int size;
} block;

void *allocate( heap *h, unsigned long long int size);
void deallocate( void * p);
```

When a data block with `nBytes` is reserved with the `alloc()` operation, it has two parts:

- Metadata:
 - o **available field** which is a long long int (8 bytes). It can have the value `AVAILABLE` (1) or `UNAVAILABLE` (0).
 - o **size field** which is a long long int (8 bytes). Contains the size in bytes of the assigned zone.
- Data: `nBytes` that the caller of the `allocate` operation can use as they see fit.

The following figure presents a block with 100 available bytes:




Figure 3 - A memory block with its descriptor

The state of the memory space is represented in the heap structure, which contains the following fields:

- *base*: initial address of the managed area
- *top*: address of the last byte allocated by a previous `alloc()` operation
- *limit*: last address of the managed area

The following figure shows an example of a scenario after performing several `alloc()` and `dealloc()` operations:




Figure 4 - Overview

Work to be developed

In this work we want to implement, in x86_64 assembly, the allocate and deallocate functions whose behavior is described below.

`void *allocate(heap *h, unsigned long long int bytesToAllocate);`

This function returns the address of a block with a size greater than or equal to *bytesToAllocate*. To achieve this, it traverses the memory described in Figure 4, starting at *h->base* and jumping to the next block structure.

The traversal ends when one of the following two conditions is met:

1. a block structure *b* is found which meets the condition $(b->available == 1) \&\& (b->size \geq bytesToAllocate)$. In this case, it returns to the user the address $b+sizeof(block)$ and sets $b->available$ to 0.
2. if it exceeds the *h->top* value. In this case, a new block with size *bytesToAllocate* must be created. If successful, the procedure is the same as in 1; the *h* structure must be updated. If it is not possible to create the new block because it would exceed the address *h->limit*, the function should do nothing and return 0.

void deallocate(void *address)

The received address is the data block (see Figure 3). The function will have to place the constant AVAILABLE (1) in the available field of the block structure.

Available files

The following files are available on CLIP:

- the myAlloc.h file that defines the API to use.
- a complete C main.c program. This program is invoked with two arguments:
./main sizeOfMemory maxBlock
where sizeOfMemory is the size of the memory managed by the allocate() and deallocate() functions. maxBlock is the maximum size of the memory block used in the tests performed. The file includes the invocation of the mmap() system call which obtains the memory area to be managed; the adjustToMultipleOfPageSize function adjusts sizeOfMemory to the smallest multiple of the page size used by the x86-64. The part of the performed tests can be modified.
- a file with a skeleton of the program in assembly called myAlloc.s. You will need to complete this file by writing the code for the allocate and deallocate functions with the behavior described above.
- a makefile that generates an executable file.

```
CC=gcc
CFLAGS=-g -z noexecsatchk
ASFLAGS=-g -gstabs
all: main
main: main.c myAlloc.o
    $(CC) $(CFLAGS) -static -o main main.c myAlloc.o
myAlloc.o: myAlloc.s
    as $(ASFLAGS) -o myAlloc.o myAlloc.s
clean:
    rm -f *.o main
```

Examples of results

As an example, a correct implementation of the main program should produce output similar to the following when executed with “./main 8192 10”:

Available	Size
1	1
1	2
1	4
1	8
Available	Size
0	1
0	2
1	4
0	8
0	10

Delivery method

The file with the x86-64 assembly functions should be named *myAlloc.s* This file should be included in an email message to be sent to master João Afonso Vilalonga's email address.

j.vilalonga@campus.fct.unl.pt

The subject of the message should be AC2024-TPC4 students XXXXX and YYYYY.

XXXXX is the student number of the 1st author and YYYYYY is the number of the 2nd author. If the group has only one member, YYYYYY must be 00000.